

modeling

October 29, 2024

0.1 Developer Salary Estimator - Modeling

Author: Topaz Montague

Deliverable Description: The goal of this Modeling Deliverable is to create a straightforward yet effective machine learning model to predict developer salaries using core factors like age, education level, years of coding experience, role type, and work experience. We'll start by splitting the data into training and testing sets, focusing on key features, and training models like Linear Regression, Decision Tree, and Random Forest to strike a good balance between accuracy and simplicity. For evaluation, we'll measure each model's performance using metrics such as Mean Absolute Error (MAE) and Mean Squared Error (MSE). The objective is to select a model that not only makes accurate predictions but also sheds light on the main drivers of developer compensation.

Project Artifacts: GitHub Repository Link: [Developer Salary GitHub Repository](#)

Overleaf Project Report: [Developer Salary Overleaf Project Report](#)

Shiny App Dashboard: [Developer Salary Shiny App Dashboard - PLACEHOLDER](#)

Data Preparation for Modeling

Load the Necessary Libraries and Data

```
[201]: # Import necessary libraries
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder, LabelEncoder

# Load the dataset
data_path = 'data/cleaned/Transformed_Developer_Survey_Data.csv'
data = pd.read_csv(data_path)

# Display the first few rows to confirm it's loaded correctly
print(data.head())
```

	ResponseId	Age	Range	RemoteWork	EdLevel	YearsCode	\
0	390	30	Remote	Some college	7		

1	399	50	Remote	Some college	38
2	417	40	Remote	Masters	21
3	427	20	Remote	Bachelors	9
4	429	30	Remote	Bachelors	20

	DevType	OrgSize	\
0	Student	15	
1	Developer, full-stack	2500	
2	Developer, back-end	250	
3	Developer, embedded applications or devices	2500	
4	Engineer, site reliability	250	

	Country	ICorPM	WorkExp	...	\
0	United States of America	Individual contributor	8	...	
1	United States of America	Individual contributor	30	...	
2	Brazil	Individual contributor	17	...	
3	Ukraine	Individual contributor	4	...	
4	United States of America	Individual contributor	15	...	

	Database_Oracle	Database_PostgreSQL	Database_Presto	Database_RavenDB	\
0	No	No	No	No	
1	No	Yes	No	No	
2	No	Yes	No	No	
3	No	No	No	No	
4	No	No	No	No	

	Database_Redis	Database_SQLite	Database_Snowflake	Database_Solr	\
0	No	Yes	No	No	
1	No	No	Yes	No	
2	No	No	No	No	
3	No	No	No	No	
4	Yes	Yes	No	No	

	Database_Supabase	Database_TiDB
0	No	No
1	No	No
2	No	No
3	No	No
4	No	No

[5 rows x 96 columns]

Prepare the Data for Modeling

```
[202]: # Step 1: Select Relevant Columns
selected_columns = ['TotalComp', 'Age Range', 'EdLevel', 'YearsCode', '
    ↪ 'DevType', 'OrgSize', 'WorkExp', 'Industry']
data = data[selected_columns]
```

```

# Step 2: Handle Missing Values

# Drop rows where TotalComp is missing (since it's the target variable)
data = data.dropna(subset=['TotalComp'])

# Handle missing values in other columns
# Check proportion of missing values
missing_values = data.isnull().mean()

# Drop or impute missing values based on proportion threshold
threshold = 0.1 # example threshold: if more than 10% of values are missing,
↳ drop the column; otherwise, drop rows
columns_to_drop = missing_values[missing_values > threshold].index
data = data.drop(columns=columns_to_drop)

# For columns with fewer missing values, drop rows with missing values
data = data.dropna()

# Step 3: Encode Categorical Variables

# List of categorical columns
categorical_columns = ['Age Range', 'EdLevel', 'DevType', 'Industry']

# Apply one-hot encoding for linear models or label encoding for tree-based
↳ models
use_one_hot = True # Set this to False if using tree-based models

if use_one_hot:
    # Use OneHotEncoder for linear models
    data = pd.get_dummies(data, columns=categorical_columns, drop_first=True)
else:
    # Use LabelEncoder for tree-based models
    label_encoders = {}
    for column in categorical_columns:
        le = LabelEncoder()
        data[column] = le.fit_transform(data[column])
        label_encoders[column] = le # Store encoder for potential inverse
↳ transformation

# Display the prepared data
print(data.head())

```

	TotalComp	YearsCode	OrgSize	WorkExp	Age Range_20	Age Range_30	\
0	110000	7	15	8	False	True	
1	195000	38	2500	30	False	False	
2	170000	21	250	17	False	False	

3	50000	9	2500	4	True	False
4	230000	20	250	15	False	True

	Age Range_40	Age Range_50	Age Range_60	Age Range_70	...	\
0	False	False	False	False	...	
1	False	True	False	False	...	
2	True	False	False	False	...	
3	False	False	False	False	...	
4	False	False	False	False	...	

	Industry_Healthcare	Industry_Higher Education	Industry_Insurance	\
0	False	False	False	
1	True	False	False	
2	False	False	False	
3	False	False	False	
4	False	False	False	

	Industry_Internet Telecomm or Information Services	Industry_Manufacturing	\
0	False	False	
1	False	False	
2	False	False	
3	False	False	
4	False	False	

	Industry_Media & Advertising Services	Industry_Other	\
0	False	False	
1	False	False	
2	False	True	
3	False	False	
4	False	False	

	Industry_Retail and Consumer Services	Industry_Software Development	\
0	False	False	
1	False	False	
2	False	False	
3	False	True	
4	False	True	

	Industry_Transportation or Supply Chain
0	True
1	False
2	False
3	False
4	False

[5 rows x 63 columns]

Get Basic Information on the Features Use `hist()` to look at the distribution of the selected features

```
[203]: # Load the dataset
df = pd.read_csv('data/cleaned/Transformed_Developer_Survey_Data.csv')

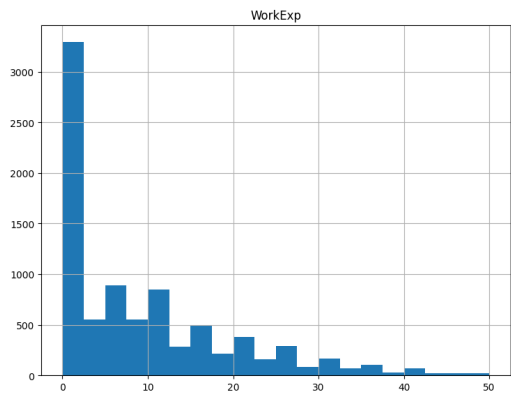
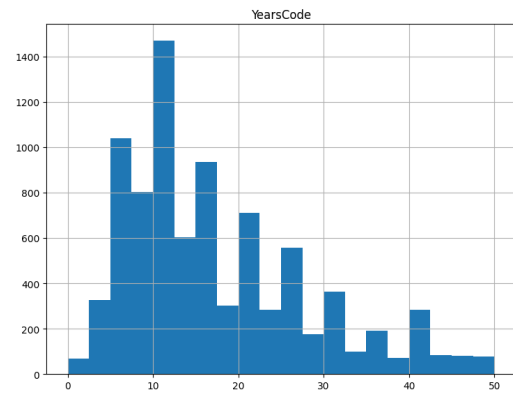
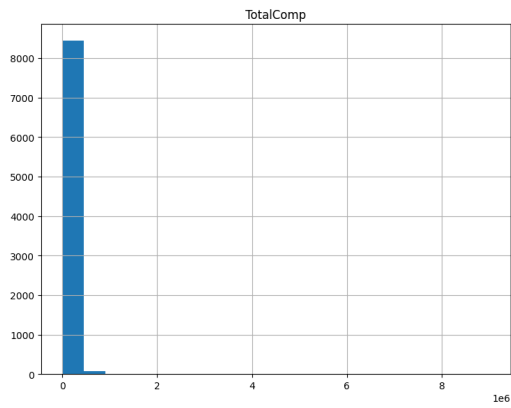
# Select only the relevant columns
selected_columns = ["TotalComp", "Age Range", "EdLevel", "YearsCode", "DevType", "OrgSize", "WorkExp", "Industry"]
df = df[selected_columns]

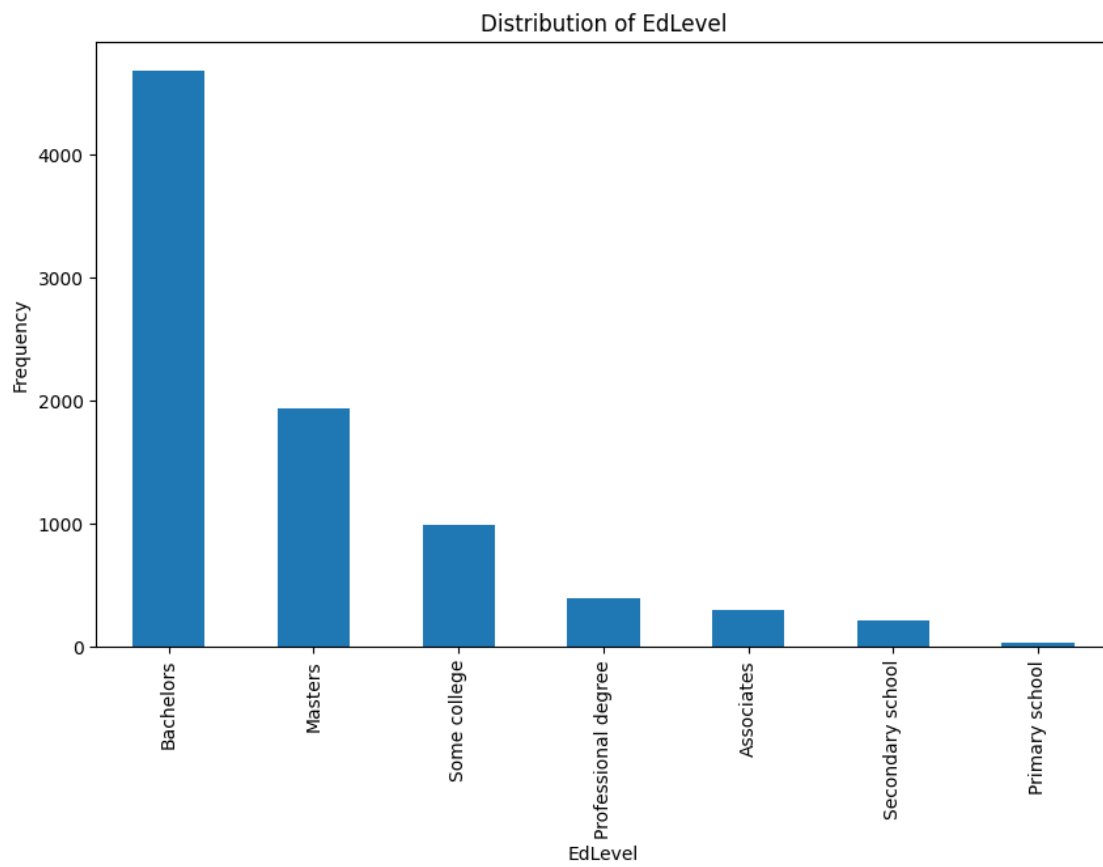

# Separate numerical and categorical columns
numerical_columns = ["TotalComp", "YearsCode", "OrgSize", "WorkExp"]
categorical_columns = ["Age Range", "EdLevel", "DevType", "Industry"]


# Plot histograms for numerical features
df[numerical_columns].hist(bins=20, figsize=(20, 15))
plt.suptitle("Distribution of Numerical Features")
plt.show()


# Plot bar charts for categorical features
for col in categorical_columns:
    plt.figure(figsize=(10, 6))
    df[col].value_counts().plot(kind='bar')
    plt.title(f"Distribution of {col}")
    plt.xlabel(col)
    plt.ylabel("Frequency")
    plt.show()
```

Distribution of Numerical Features









```
[204]: # Check that cleaning is good
df.head(n=10)
```

```
[204]:
```

	TotalComp	Age Range	EdLevel	YearsCode \
0	110000	30	Some college	7
1	195000	50	Some college	38
2	170000	40	Masters	21
3	50000	20	Bachelors	9
4	230000	30	Bachelors	20
5	85000	40	Bachelors	25
6	160000	40	Secondary school	20
7	110000	60	Bachelors	25
8	190000	40	Masters	23
9	115000	50	Associates	10

	DevType	OrgSize	WorkExp	\
0	Student	15	8	
1	Developer, full-stack	2500	30	
2	Developer, back-end	250	17	
3	Developer, embedded applications or devices	2500	4	
4	Engineer, site reliability	250	15	
5	Developer, full-stack	50	25	
6	Developer, full-stack	15	20	
7	Developer, front-end	2500	0	
8	Engineering manager	10000	0	
9	Developer, full-stack	250	10	

	Industry
0	Transportation or Supply Chain
1	Healthcare
2	Other
3	Software Development
4	Software Development
5	Retail and Consumer Services
6	Software Development
7	Other
8	Other
9	Other

Plot to Visualize the Selected Features vs Total Compensation Total Compensation vs Age Range

```
[205]: import pandas as pd
import matplotlib.pyplot as plt

# Load the dataset
file_path = 'data/cleaned/Transformed_Developer_Survey_Data.csv'
data = pd.read_csv(file_path)

# Filter data to include only rows where TotalComp <= 400,000
data = data[data['TotalComp'] <= 400000]

# Define color mapping (specific to the column you're working with)
color_map = {
    20: 'green', 30: 'blue', 40: 'purple', 50: 'orange', 60: 'red', 70: 'cyan'
}

# Plotting example for TotalComp vs Age Range
plt.figure(figsize=(10, 6))

# Loop through unique age ranges to create separate scatter plots
```

```

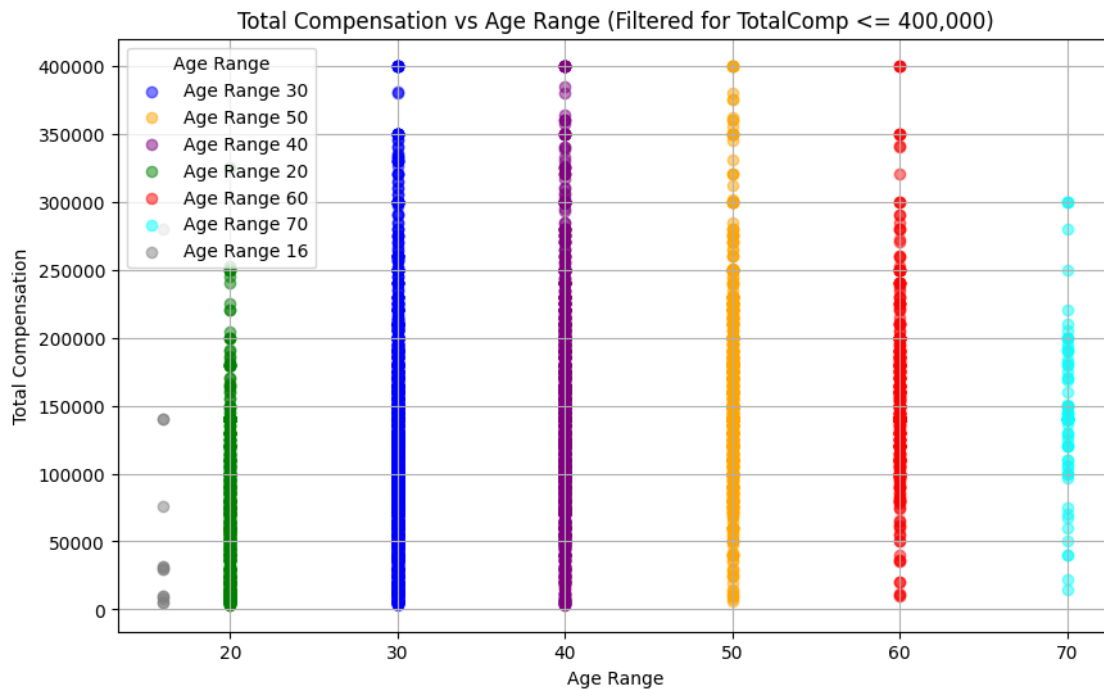
for age_range in data['Age Range'].unique():
    subset = data[data['Age Range'] == age_range]
    plt.scatter(
        subset['Age Range'],
        subset['TotalComp'],
        c=color_map.get(age_range, 'grey'), # Default to 'grey' if age_range
        ↪not in color_map
        alpha=0.5,
        label=f'Age Range {age_range}'
    )

plt.title('Total Compensation vs Age Range (Filtered for TotalComp <= 400,000)')
plt.xlabel('Age Range')
plt.ylabel('Total Compensation')

# Set y-axis to display values in plain format (no scientific notation)
plt.ticklabel_format(style='plain', axis='y')

plt.legend(title='Age Range')
plt.grid(True)
plt.show()

```

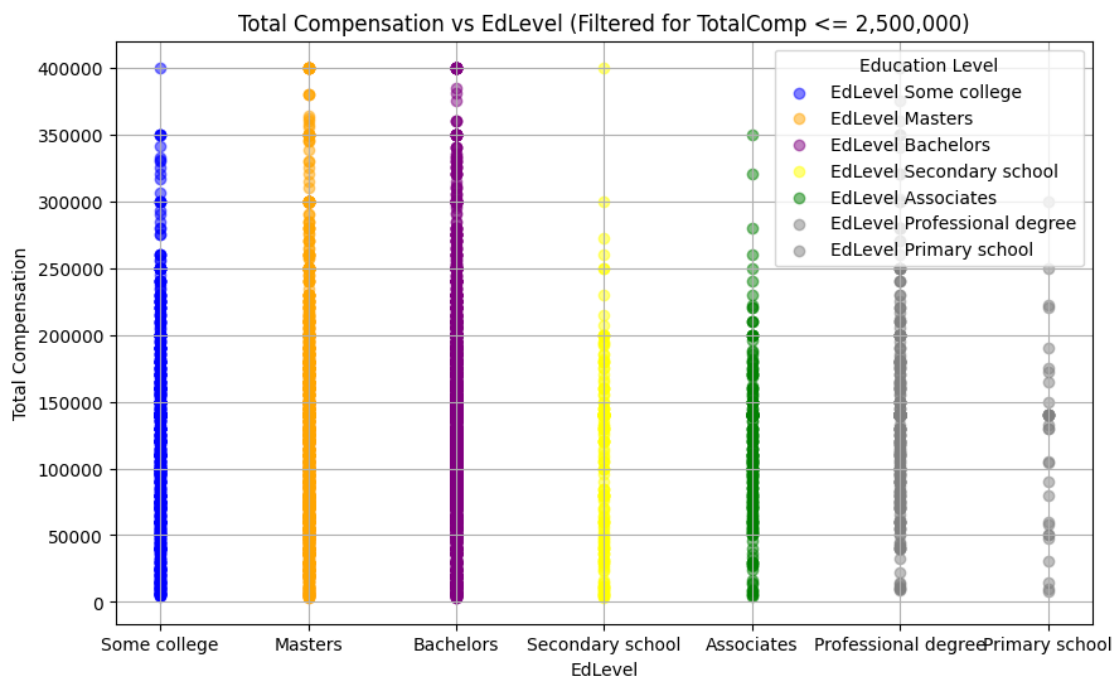


Total Compensation vs Education Level

```
[206]: # Define color mapping for EdLevel
color_map = {
    'High school': 'green', 'Some college': 'blue', 'Bachelors': 'purple',
    'Masters': 'orange', 'Doctorate': 'red', 'Secondary school': 'yellow',
    'Associates': 'green', 'Professional Degree': 'purple'
}

plt.figure(figsize=(10, 6))
for ed_level in data['EdLevel'].unique():
    subset = data[data['EdLevel'] == ed_level]
    plt.scatter(
        subset['EdLevel'],
        subset['TotalComp'],
        c=color_map.get(ed_level, 'grey'),
        alpha=0.5,
        label=f'EdLevel {ed_level}'
    )

plt.title('Total Compensation vs EdLevel (Filtered for TotalComp <= 2,500,000)')
plt.xlabel('EdLevel')
plt.ylabel('Total Compensation')
plt.ticklabel_format(style='plain', axis='y')
plt.legend(title='Education Level')
plt.grid(True)
plt.show()
```



Total Compensation vs Years of Coding Experience

```
[207]: import pandas as pd
import matplotlib.pyplot as plt
import matplotlib.cm as cm

# Load the dataset
file_path = 'data/cleaned/Transformed_Developer_Survey_Data.csv'
data = pd.read_csv(file_path)

# Filter dataset to only include TotalComp <= 400,000
filtered_data = data[data['TotalComp'] <= 400000]

# Define figure and axis
fig, ax = plt.subplots(figsize=(10, 6))

# Normalize the YearsCode data for the color map
norm = plt.Normalize(filtered_data['YearsCode'].min(),
    ↪filtered_data['YearsCode'].max())
colors = cm.Blues(norm(filtered_data['YearsCode']))

# Plotting TotalComp vs YearsCode with a gradient blue color based on YearsCode
scatter = ax.scatter(
    filtered_data['YearsCode'],
    filtered_data['TotalComp'],
    c=filtered_data['YearsCode'],
    cmap='Blues',
    alpha=0.7
)


# Add color bar to indicate the gradient scale of YearsCode
cbar = fig.colorbar(scatter, ax=ax)
cbar.set_label('Years of Coding Experience')

# Set titles and labels
ax.set_title('Total Compensation vs YearsCode (Filtered for TotalComp <=
    ↪400,000)')
ax.set_xlabel('Years of Coding Experience')
ax.set_ylabel('Total Compensation')

# Display y-axis in plain format without scientific notation
ax.ticklabel_format(style='plain', axis='y')

# Add grid for readability
ax.grid(True, linestyle='--', alpha=0.7)

plt.show()
```



Total Compensation vs Type of Developer

```
[208]: import pandas as pd
import matplotlib.pyplot as plt

# Load the dataset
file_path = 'data/cleaned/Transformed_Developer_Survey_Data.csv'
data = pd.read_csv(file_path)

# Filter dataset to only include TotalComp <= 400,000
filtered_data = data[data['TotalComp'] <= 400000]

# Define figure with adjusted width
plt.figure(figsize=(15, 8))

# Plotting TotalComp vs DevType with appropriate adjustments for readability
for dev_type in filtered_data['DevType'].unique():
    subset = filtered_data[filtered_data['DevType'] == dev_type]
    plt.scatter(
        subset['DevType'],
        subset['TotalComp'],
        alpha=0.5,
        label=f'DevType {dev_type}'
    )
```

```

# Title and labels
plt.title('Total Compensation vs DevType (Filtered for TotalComp <= 400,000)')
plt.xlabel('Developer Type')
plt.ylabel('Total Compensation')

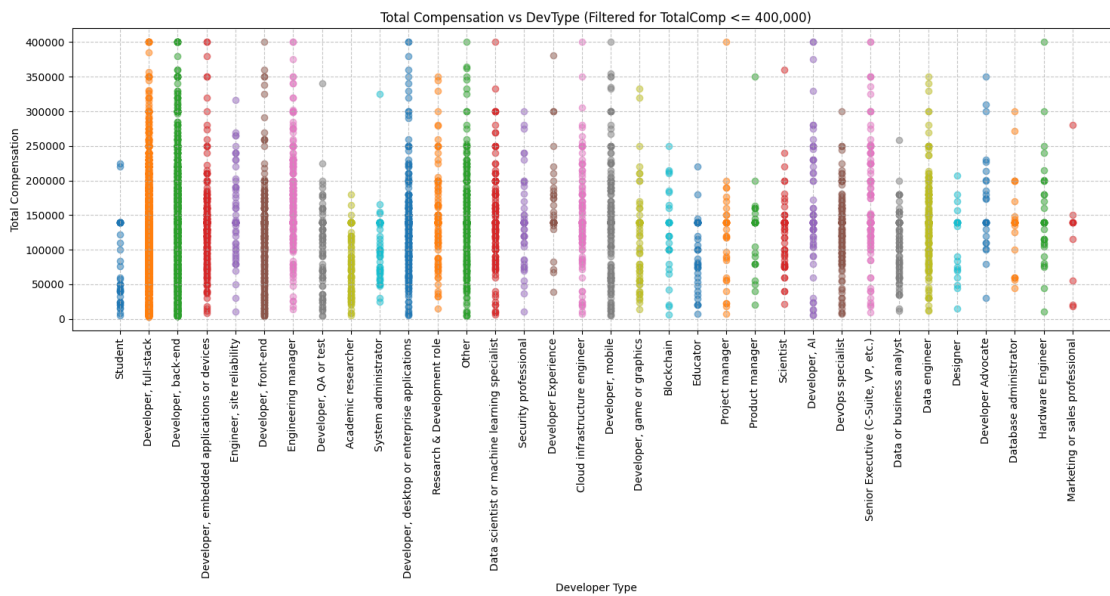
# Display y-axis in plain format without scientific notation
plt.ticklabel_format(style='plain', axis='y')

# Rotate x-axis labels for readability
plt.xticks(rotation=90)

# Add grid for readability
plt.grid(True, linestyle='--', alpha=0.7)

plt.tight_layout() # Adjust layout to fit everything neatly
plt.show()

```



Total Compensation vs Organization Size

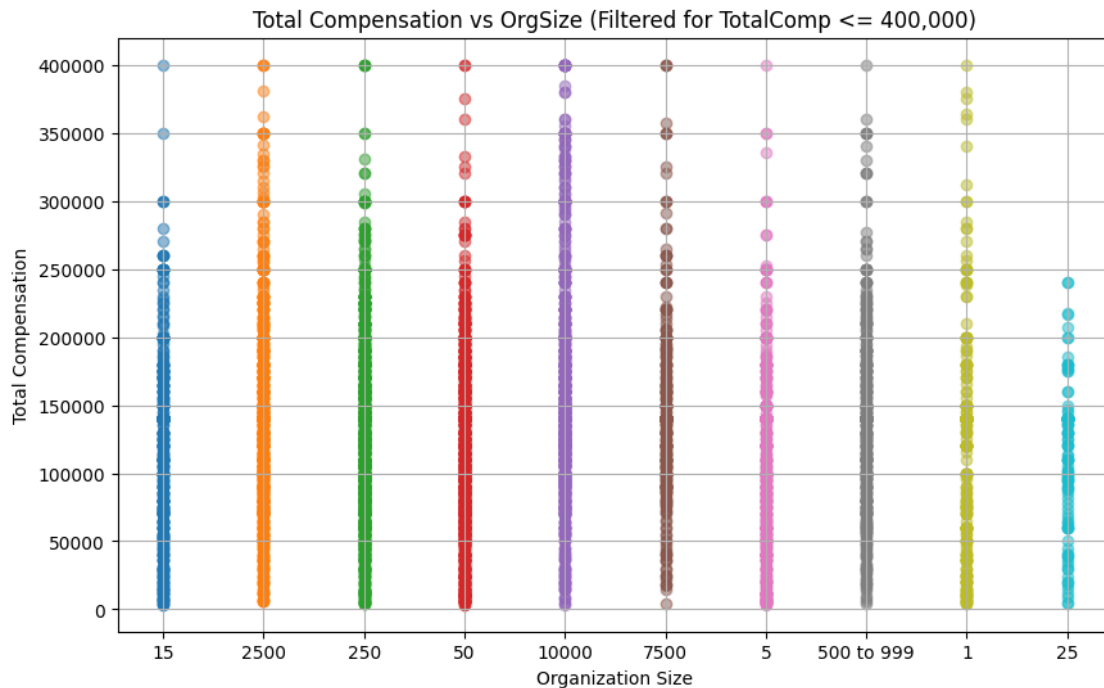
```

[209]: plt.figure(figsize=(10, 6))
for org_size in filtered_data['OrgSize'].unique():
    subset = filtered_data[filtered_data['OrgSize'] == org_size]
    plt.scatter(
        subset['OrgSize'],
        subset['TotalComp'],
        alpha=0.5,
        label=f'OrgSize {org_size}'
    )

```

```
)

plt.title('Total Compensation vs OrgSize (Filtered for TotalComp <= 400,000)')
plt.xlabel('Organization Size')
plt.ylabel('Total Compensation')
plt.ticklabel_format(style='plain', axis='y')
plt.grid(True)
plt.show()
```



Declare the Train-Test Split

```
[210]: # Use train_test_split from sklearn.model_selection to split the data into
        ↪ training and testing sets. A typical split is 80% for training and 20% for
        ↪ testing:
from sklearn.model_selection import train_test_split

# Filter data to include only rows where TotalComp <= 400000
filtered_data = data[data['TotalComp'] <= 400000]

# Define features and target
X = filtered_data[['Age Range', 'EdLevel', 'YearsCode', 'DevType', 'OrgSize',
        ↪ 'WorkExp', 'Industry']]
y = filtered_data['TotalComp']

# Split the data into training and testing sets
```



```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)
```

Model Training and Evaluation Linear Regression

```
[211]: import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

# Load the dataset
data = pd.read_csv('data/cleaned/Transformed_Developer_Survey_Data.csv')

# Filter data to include only rows where TotalComp <= 400000
data = data[data['TotalComp'] <= 400000]

# Select relevant columns
selected_columns = ['TotalComp', 'Age Range', 'EdLevel', 'YearsCode',
↳'DevType', 'OrgSize', 'WorkExp', 'Industry']
data = data[selected_columns]

# Handle missing values (drop rows with missing values for simplicity)
data = data.dropna()

# One-hot encode categorical variables (including 'OrgSize')
data = pd.get_dummies(data, columns=['Age Range', 'EdLevel', 'DevType',
↳'OrgSize', 'Industry'], drop_first=True)

# Define features and target
X = data.drop('TotalComp', axis=1)
y = data['TotalComp']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)

# Initialize and train the model
model = LinearRegression()
model.fit(X_train, y_train)

# Predict on test data
y_pred = model.predict(X_test)

# Evaluate the model
print("MAE:", mean_absolute_error(y_test, y_pred))
print("MSE:", mean_squared_error(y_test, y_pred))
print("R2 Score:", r2_score(y_test, y_pred))
```

MAE: 41487.11481678698
MSE: 3243122526.822035
R2 Score: 0.19654005605061176

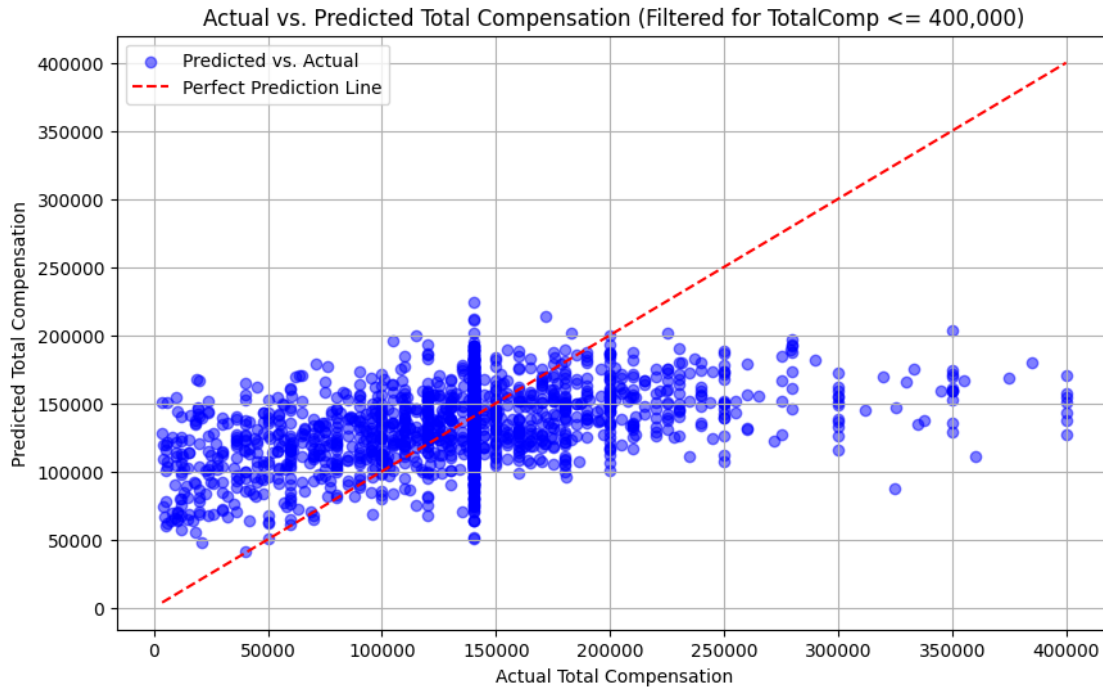
Linear Regression Model Graph for Total Compensation Predictions vs Actuals

```
[212]: import matplotlib.pyplot as plt

# Filter the data for TotalComp <= 400000
y_test_filtered = y_test[y_test <= 400000]
y_pred_filtered = y_pred[y_test <= 400000] # Apply the same mask to y_pred

# Scatter plot of actual vs predicted values
plt.figure(figsize=(10, 6))
plt.scatter(y_test_filtered, y_pred_filtered, alpha=0.5, color='blue',
            label='Predicted vs. Actual')
plt.plot([y_test_filtered.min(), y_test_filtered.max()],
         [y_test_filtered.min(), y_test_filtered.max()],
         color='red', linestyle='--', label='Perfect Prediction Line')

# Labeling the plot
plt.xlabel("Actual Total Compensation")
plt.ylabel("Predicted Total Compensation")
plt.title("Actual vs. Predicted Total Compensation (Filtered for TotalComp <= 400,000)")
plt.legend()
plt.grid(True)
plt.ticklabel_format(style='plain', axis='y') # Avoid scientific notation on the y-axis
plt.show()
```



Decision Tree Regressor

```
[213]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

# Load the dataset
file_path = 'data/cleaned/Transformed_Developer_Survey_Data.csv'
data = pd.read_csv(file_path)

# Filter data to include only rows where TotalComp <= 400000
data = data[data['TotalComp'] <= 400000]

# Select relevant columns
selected_columns = ['TotalComp', 'Age Range', 'EdLevel', 'YearsCode', '
    ↳ DevType', 'OrgSize', 'WorkExp', 'Industry']
data = data[selected_columns]

# Handle missing values
data = data.dropna(subset=['TotalComp']) # Drop rows with missing TotalComp
    ↳ values
data = data.dropna() # Drop rows with missing values in other columns

# Encode categorical variables
```

```

categorical_columns = ['Age Range', 'EdLevel', 'DevType', 'OrgSize', 'Industry']
data = pd.get_dummies(data, columns=categorical_columns, drop_first=True)

# Define features and target
X = data.drop('TotalComp', axis=1)
y = data['TotalComp']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

# Initialize and train the Decision Tree Regressor model
model = DecisionTreeRegressor(random_state=42)
model.fit(X_train, y_train)

# Predict on test data
y_pred = model.predict(X_test)

# Evaluate the model
print("MAE:", mean_absolute_error(y_test, y_pred))
print("MSE:", mean_squared_error(y_test, y_pred))
print("R2 Score:", r2_score(y_test, y_pred))

```

MAE: 60327.253656967536

MSE: 6833384163.111126

R2 Score: -0.6929210695156642

Decision Tree Regressor Model Graph for Total Compensation Predictions vs Actuals

```

[214]: # Visualize Actual vs Predicted TotalComp for values <= 400000
import matplotlib.pyplot as plt


# Filter the data for TotalComp <= 400000
y_test_filtered = y_test[y_test <= 400000]
y_pred_filtered = y_pred[y_test <= 400000] # Apply the same mask to y_pred

# Plot actual vs predicted values
plt.figure(figsize=(10, 6))
plt.scatter(y_test_filtered, y_pred_filtered, alpha=0.5, color='blue',
    label='Predicted vs. Actual')
plt.plot([y_test_filtered.min(), y_test_filtered.max()], [y_test_filtered.
    min(), y_test_filtered.max()],
    color='red', linestyle='--', label='Perfect Prediction Line')

# Labeling the plot
plt.xlabel("Actual Total Compensation")
plt.ylabel("Predicted Total Compensation")

```

```
plt.title("Decision Tree Regressor: Actual vs. Predicted Total Compensation_␣
↳(Filtered for TotalComp <= 400000)")
plt.legend()
plt.grid(True)
plt.ticklabel_format(style='plain')
plt.show()
```



Random Forest Regressor

```
[215]: # Import necessary libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

# Load the dataset
file_path = 'data/cleaned/Transformed_Developer_Survey_Data.csv'
data = pd.read_csv(file_path)

# Filter data to include only rows where TotalComp <= 400000
data = data[data['TotalComp'] <= 400000]

# Select relevant columns
selected_columns = ['TotalComp', 'Age Range', 'EdLevel', 'YearsCode', ␣
↳'DevType', 'OrgSize', 'WorkExp', 'Industry']
```

```

data = data[selected_columns]

# Handle missing values (drop rows with missing values for simplicity)
data = data.dropna()

# Encode categorical variables
categorical_columns = ['Age Range', 'EdLevel', 'DevType', 'OrgSize', 'Industry']
data = pd.get_dummies(data, columns=categorical_columns, drop_first=True)

# Define features and target
X = data.drop('TotalComp', axis=1)
y = data['TotalComp']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=42)

# Initialize and train the Random Forest Regressor model
model = RandomForestRegressor(random_state=42)
model.fit(X_train, y_train)

# Predict on test data
y_pred = model.predict(X_test)

# Evaluate the model
print("MAE:", mean_absolute_error(y_test, y_pred))
print("MSE:", mean_squared_error(y_test, y_pred))
print("R2 Score:", r2_score(y_test, y_pred))

```

MAE: 45342.22948174411

MSE: 3782466515.6126447

R2 Score: 0.06292151792285938

Random Forest Regressor Model Graph for Total Compensation Predictions vs Actuals

```

[216]: import matplotlib.pyplot as plt

# Filter to only include TotalComp values <= 400,000
y_test_filtered = y_test[y_test <= 400000]
y_pred_filtered = y_pred[y_test <= 400000] # Apply the same mask to y_pred

# Visualize Actual vs Predicted TotalComp (Filtered for <= 400,000)
plt.figure(figsize=(10, 6))

# Scatter plot of actual vs predicted values
plt.scatter(y_test_filtered, y_pred_filtered, alpha=0.5, color='blue',
            label='Predicted vs. Actual')

```

```

# Plot a perfect prediction line
plt.plot([y_test_filtered.min(), y_test_filtered.max()], [y_test_filtered.
    ↪min(), y_test_filtered.max()],
    color='red', linestyle='--', label='Perfect Prediction Line')

# Labeling the plot
plt.xlabel("Actual Total Compensation")
plt.ylabel("Predicted Total Compensation")
plt.title("Random Forest Regressor: Actual vs. Predicted Total Compensation_
    ↪(Filtered for TotalComp <= 400,000)")
plt.legend()
plt.grid(True)

# Set y-axis to display values in plain format (no scientific notation)
plt.ticklabel_format(style='plain', axis='both')
plt.show()

```

